WINDOWS THREAT HUNTING WITH HAYABUSA

Overnight execution plan for a team of five

Working window: 8:00 PM – 8:00 AM (12 hours)

Target by morning	A reproducible Windows Event Log analysis demo, a small set of validated findings, an ATT&CK-mapped report, and a short presentation.
Project principle	Use Hayabusa as the detection/log-analysis engine. The team's contribution is the investigation workflow, evidence validation, mapping, reporting, and optional lightweight interface.
Safety boundary	Use a lab VM and synthetic, sample, or explicitly authorized Windows event logs. Do not investigate devices or logs without permission. Treat alerts as leads, not proof of compromise.

This plan is designed to prevent five people from independently producing five incompatible projects at 5 AM.

1. Scope and definition of done

Project title: **Windows Threat Hunting System Using Hayabusa**

The team will parse Windows Event Logs with Hayabusa, review prioritized detections, validate a small number of findings against the underlying event evidence, map supported findings to MITRE ATT&CK where appropriate, and produce a repeatable report. A simple dashboard is optional; a working, documented command-line workflow is more valuable than a half-finished interface.

By 8:00 AM, the team should have:

- A documented lab setup and a clearly identified, authorized log dataset.
- Hayabusa installed and successfully run against the chosen dataset, with the exact command and version recorded.
- Raw input logs preserved separately from generated output; no evidence silently edited.
- At least 3 candidate detections reviewed, with 1–3 findings documented in depth if the dataset supports them. If it does not, report that honestly rather than inventing incidents.
- A findings table with timestamps, event/log references, rule or alert name, severity, interpretation, confidence, limitations, and ATT&CK mapping only when justified.
- A final report, a 5–7 slide presentation, a demo checklist, and a shared folder with clearly named files.
- A reproducible runbook so another teammate can repeat the analysis.

Important distinction: a rule match is not automatically a confirmed attack. Label results as **candidate / validated observation / inconclusive** and explain what evidence supports the label.

2. Team roles and ownership

Member 1 — Project Lead & Integration Owner

Role purpose: Owns scope, schedule, coordination, and final consistency. This person is not the team's full-time messenger; they also own the project narrative and integration.

Responsibilities

- At 8–8:30 PM, create the shared folder and task board; confirm the dataset, lab boundaries, and expected deliverables.
- Maintain a single source-of-truth README with current status, file paths, owners, and decisions.
- Set checkpoints, resolve overlaps, and stop scope creep. Keep the optional dashboard from blocking the core analysis.
- Collect outputs from Members 2–5, standardize terminology and filenames, and check that report claims match evidence.
- Own the final report structure, presentation flow, demo order, and final quality gate.

Must submit to the team

- Project README and task/status tracker.
- Integrated final report (PDF or DOCX), presentation deck, and demo script/checklist.
- Final evidence-to-claim review checklist and a short list of known limitations.

Handoff / owned work: Project folder structure, README, timeline, final report, slide deck, demo checklist.

Member 2 — Lab Setup & Log Acquisition

Role purpose: Owns the authorized Windows lab and input dataset. Ensures the team can reproduce where logs came from and how they were prepared.

Responsibilities

- Confirm use of a Windows test VM or approved sample dataset; document OS version, time zone, and dataset source.
- Install/verify prerequisites needed by the chosen Hayabusa release; record version numbers and setup steps.
- Identify relevant EVTX files or export a permitted test dataset. Keep originals read-only and make a separate working copy.

- Create a dataset inventory: filename, approximate time coverage, source host/context if known, and any gaps.
- Run a quick sanity check that files are readable and timestamps/time zone assumptions are documented.
- Hand off the exact input directory and setup/run notes to Member 3; remain available if parsing fails.

Must submit to the team

- lab_setup.md with environment, installation notes, version, and troubleshooting.
- dataset_inventory.csv or a clear inventory table.
- Input dataset folder or source reference, plus a note confirming authorization and any constraints.

Handoff / owned work: Lab environment, data inventory, setup instructions, preserved input logs.

Member 3 — Hayabusa Analysis Engineer

Role purpose: Owns the core tool execution and the first-pass alert output. This is the main technical engine room.

Responsibilities

- Verify the Hayabusa command syntax against the installed version's help/documentation rather than copying an outdated command.
- Run Hayabusa against the agreed dataset using a documented output format and destination.
- Save the exact command, tool version, run time, output filenames, and any warnings/errors.
- Review output for high-priority or explainable candidate detections; select a manageable shortlist for manual review.
- For each shortlisted alert, capture rule/alert name, timestamp, relevant event IDs/fields, severity, and output row or source reference.
- Rerun the pipeline once after corrections to confirm the runbook works from a clean output directory.

Must submit to the team

- hayabusa_runbook.md with exact commands and version information.
- Unmodified raw Hayabusa output (CSV/JSON or the supported format selected by the team).
- shortlist.csv containing candidate detections and evidence pointers.
- A brief note of errors, false positives, and parsing limitations.

Handoff / owned work: Hayabusa execution, reproducible commands, raw results, initial alert shortlist.

Member 4 — Threat Hunter & ATT&CK; Mapper

Role purpose: Owns interpretation and evidence validation. Prevents the project from treating every alert as a confirmed breach, a common but avoidable academic tradition.

Responsibilities

- Take the alert shortlist from Member 3 and inspect the supporting event fields or related events.
- Build a small timeline for selected candidates; normalize timestamps and explicitly note the time zone.
- For each candidate, explain what behavior may be indicated, what evidence supports it, what benign explanation is possible, and what additional evidence would be needed.
- Assign a confidence/status label: candidate, validated observation, or inconclusive. Do not claim maliciousness without supporting evidence.
- Map behavior to MITRE ATT&CK; technique/sub-technique only where the evidence supports the behavior; include technique ID and a one-sentence rationale.
- If the dataset has no strong suspicious behavior, document a defensible benign/limited result and describe what the dataset cannot show.

Must submit to the team

- findings.csv or findings.md with evidence, interpretation, confidence, and limitations.
- timeline.csv or a compact chronological table.
- ATT&CK; mapping table with technique IDs, rationale, and evidence references.
- One paragraph per selected finding suitable for the final report.

Handoff / owned work: Validated findings, timeline, ATT&CK; mapping, confidence and limitation statements.

Member 5 — Reporting, Interface & QA Owner

Role purpose: Owns how results are presented and tested. Builds only a lightweight interface if the core pipeline is already working.

Responsibilities

- Agree on the report schema with Members 1, 3, and 4: timestamp, alert/rule, severity, event reference, interpretation, status/confidence, ATT&CK; ID, recommended next step.
- Prepare a clean summary table or basic local dashboard/HTML page that displays the selected findings; do not spend the night building a full product.
- Capture genuine screenshots of the tool run and output; redact host/user details if they are not necessary for the demo.
- Check file paths, timestamps, labels, broken links, spelling, consistent severity terminology, and whether each claim has an evidence reference.
- Run the demo checklist from a fresh terminal/session or documented clean state and record what works and what does not.
- Prepare the limitations, evaluation criteria, and future improvements section with Member 1.

Must submit to the team

- summary.html, a simple table, or dashboard source (optional and time-boxed).
- Screenshots of real tool output and final results.
- qa_checklist.md with pass/fail results and known issues.
- A concise evaluation/limitations section and suggested next steps.

Handoff / owned work: Readable result summary, screenshots, QA checks, optional minimal interface.

3. Overnight schedule: 8:00 PM – 8:00 AM

The time blocks below are checkpoints, not permission to disappear into a side quest. Each member works in parallel where possible, but dependencies are explicit.

Time	All-team checkpoint	Member 1	Member 2	Member 3	Member 4	Member 5
8:00–8:30 PM	Kickoff; agree scope, data, deliverables, folder structure.	Lead kickoff; create tracker and shared folders.	Confirm authorized dataset and lab access.	Check Hayabusa availability/version requirements.	Prepare finding/evidence template and ATT&CK; reference format.	Create report outline, QA checklist, and screenshot plan.
8:30–9:30 PM	Environment and data confirmed; core command plan agreed.	Lock scope and define naming conventions.	Set up lab; inventory EVTX files; document source/time zone.	Install/verify tool; inspect help and valid command syntax.	Define what counts as evidence and confidence labels.	Build report skeleton and findings table template.
9:30–10:30 PM	First successful run or explicit blocker escalated.	Track blockers; make quick decisions.	Hand off dataset and inventory to Member 3.	Run Hayabusa; save command, version, output, errors.	Review initial output structure and shortlist candidates with Member 3.	Prepare demo folder and capture setup evidence.
10:30–11:30 PM	Initial candidate list selected; no unsupported claims.	Review progress; draft project overview/objectives.	Help resolve data gaps or test a second approved sample if needed.	Cleanly rerun if needed; export candidate shortlist.	Start evidence review and timeline for top candidates.	Build summary-table or optional simple HTML view from agreed schema.
11:30 PM–12:30 AM	Evidence validation underway; midpoint check.	Draft methodology and system workflow sections.	Document environment, dataset provenance, and limitations.	Answer rule/output questions; provide exact event references.	Validate candidate events; assess benign explanations; draft ATT&CK; mappings.	Populate report with verified tool outputs and screenshots.
12:30–1:00 AM	Short break and handoff sync.	Run 10-minute sync; log unresolved issues.	Confirm input data and setup notes are saved.	Ensure raw outputs are backed up and readable.	Share provisional finding statuses and uncertainty.	Back up current report and collect filenames.
1:00–2:00 AM	First complete findings table and timeline.	Integrate methods and results; check consistency.	Test reproducibility on a small subset or clean output folder.	Resolve any parsing/runbook gaps; freeze raw output.	Finish evidence references, timeline, ATT&CK; rationale.	Finish summary view and screenshots; don't expand UI scope.
2:00–3:00 AM	Core pipeline frozen; no feature creep.	Review the complete end-to-end story.	Finish lab_setup.md and dataset inventory.	Finish hayabusa_runbook.md and shortlist.csv.	Finish findings/timeline/ATT&CK; files.	Run QA pass on outputs and file naming.
3:00–4:00 AM	All individual deliverables submitted to shared folder.	Assemble report draft and slides.	Verify setup instructions are clear to another teammate.	Repeat documented command and confirm output exists.	Cross-check every finding against event evidence.	Run demo from clean state; record pass/fail.
4:00–5:00 AM	Integrated draft reviewed by at least two people.	Edit report for clarity and claim-evidence alignment.	Answer setup questions; confirm dataset provenance wording.	Explain tool behavior/output fields for report.	Check ATT&CK; mappings and confidence labels.	Check screenshots, tables, formatting, spelling, and links.
5:00–6:00 AM	Rehearsal version ready; only critical fixes allowed.	Lead a timed demo rehearsal; assign speaking parts.	Prepare environment for reliable live demo or fallback screenshots.	Run the exact final analysis command once more if safe and needed.	Prepare 1–2 evidence walkthroughs with caveats.	Finalize slides, QA checklist, and backup demo artifacts.
6:00–7:00 AM	Final validation and presentation rehearsal.	Check rubric/requirements; freeze scope.	Confirm dataset and lab can be reopened; no missing files.	Confirm raw outputs and runbook match the final run.	Verify findings, timeline, ATT&CK; IDs, limitations.	Export final PDF/deck; verify files open correctly.

Time	All-team checkpoint	Member 1	Member 2	Member 3	Member 4	Member 5
7:00–7:30 AM	Final handoff and backup.	Approve final folder and final report version.	Archive input inventory and setup notes.	Archive raw outputs and runbook.	Archive findings and mapping table.	Archive screenshots, QA checklist, slide deck.
7:30–8:00 AM	Presentation / submission buffer.	Present overview and methodology; manage transitions.	Explain lab and dataset provenance.	Demonstrate Hayabusa execution and output.	Explain findings, evidence, and ATT&CK; mapping.	Show results summary, QA, limitations, and future work.

Checkpoint rule: if the tool has not run successfully by 10:30 PM, pause optional interface work. Member 3 and Member 2 troubleshoot together; Member 1 chooses a smaller known-good authorized sample if appropriate. The project must still produce a reproducible method and honest limitations.

4. Required shared-folder structure

Use one agreed folder structure so nobody spends the last hour asking which file is final.

```
Hayabusa-Threat-Hunting/
■■■ README.md
■■■ 01_lab-and-data/
■ ■■■ lab_setup.md
■ ■■■ dataset_inventory.csv
■■■ 02_raw-input/ (preserve originals; restrict access as needed)
■■■ 03_hayabusa-output/
■ ■■■ hayabusa_runbook.md
■ ■■■ raw_output.csv (or chosen supported format)
■ ■■■ shortlist.csv
■■■ 04-investigation/
■ ■■■ findings.csv
■ ■■■ timeline.csv
■ ■■■ attack_mapping.csv
■■■ 05-summary-view/
■ ■■■ summary.html (optional)
■ ■■■ screenshots/
■■■ 06-report-and-slides/
■ ■■■ final_report.pdf
■ ■■■ presentation.pptx
■■■ 07-qa/
■■■ qa_checklist.md
■■■ demo_checklist.md
```

5. Findings record: minimum required fields

Field	What to record	Owner
Finding ID	Stable identifier such as F-01; use it consistently in report, slides, and screenshots.	Member 4
Timestamp / time zone	Event time and the time zone assumption; do not silently mix local time and UTC.	Member 4
Alert / rule	Hayabusa alert or rule name, severity as output, and output row/reference.	Member 3
Evidence reference	Relevant event ID, log name, fields, file, and line/row/reference where possible.	Members 3 + 4
Interpretation	What the event may indicate and why it matters.	Member 4
Status / confidence	Candidate, validated observation, or inconclusive, with a short reason.	Member 4
ATT&CK; mapping	Technique ID/name and evidence-based rationale; leave unmapped if not supported.	Member 4
Limitations / next step	Missing context, plausible benign explanation, and what extra telemetry would help.	Members 4 + 5

6. Final quality gate

- The report's commands match the version of Hayabusa actually used.
- The input dataset source and authorization are documented; raw inputs are preserved.
- The analysis can be repeated from the runbook without relying on someone's memory.
- Every reported finding points to evidence. No fabricated alerts, screenshots, event fields, or results.
- ATT&CK; mappings are justified by observed behavior, not attached merely to make the report look more advanced.
- The team distinguishes tool output from analyst interpretation and states uncertainty.

- The demo has a fallback: saved output and screenshots in case the live environment fails.
- The final PDF and slides open correctly and are saved in the shared folder.

7. Scope control and contingency plan

Problem	Immediate response
Hayabusa installation or command failure	Check the installed version's help and official documentation; Member 2 checks prerequisites while Member 3 isolates the error. Use an approved known-good sample if needed.
Few or no suspicious detections	Do not invent a threat. Demonstrate the pipeline, discuss the sample's coverage, report benign/limited results, and explain what extra telemetry would be needed.
Too many alerts to review	Narrow to a small, explainable shortlist based on severity, relevance, and evidence availability. Record the selection criteria.
Dashboard unfinished	Stop UI work and use a clear CSV/table plus screenshots. The project is threat hunting, not a competition to build the fanciest rectangle.
Conflicting conclusions	Return to source event evidence, label uncertainty, and use the more conservative claim. Member 1 resolves wording only after Member 4 explains the evidence.
A teammate falls behind	Protect core dependencies first: data → Hayabusa run → evidence validation → report. Reassign polish tasks; do not duplicate core analysis blindly.

8. Communication rules for the night

- Use one team chat channel and one shared folder. Every deliverable has one named owner.
- At each checkpoint, each member reports: **Done**, **Next**, **Blocked by**, and the exact file path.
- Save work at every handoff. File names should include content, not vague labels such as final_final_latest2.
- No one overwrites raw input or raw Hayabusa output. Work from copies and preserve provenance.
- Take a 5–10 minute break at the midpoint, hydrate, and rotate away from screens briefly. Exhaustion is not a testing methodology.

9. Suggested presentation split

Speaker	Time	Content
Member 1	1 minute	Problem statement, objective, scope, and overall workflow.
Member 2	45 seconds	Lab setup, data source, and safe handling.
Member 3	1–2 minutes	Hayabusa run, key output, and reproducibility.
Member 4	2 minutes	One or two evidence-backed findings, timeline, ATT&CK; mapping, and uncertainty.
Member 5	1 minute	Results summary, QA, limitations, and future improvements.

Final reminder: A small, reproducible investigation with traceable evidence is stronger than a large feature list that does not work. Keep the objective narrow, submit the required files, and leave the team with a demo that can survive a skeptical question.